

# Spatial Statistics

PM569, Spatial Statistics

Lecture 1: August 28th, 2015

## Course Details

- ▶ Meeting time/place: Fridays 10a-1p, Soto 117
- ▶ Last class: December 4 (We will have project presentations that day)
- ▶ Prerequisites: PM 510, 511a (or equivalent intro statistics and regression modeling)
- ▶ Grading: 6 assignments (10% each), 1 final project/presentation small group or individual (30%)
- ▶ Grading: Weekly reading assignments and discussion (10%)

## Course Details

- ▶ **Assignments** must be submitted individually but you can discuss with others. No copying! Late penalty of 20% per day.
- ▶ The **final project** will also be done individually. There are several components to the project: deciding a suitable topic (making sure you have some data), writing a brief proposal, making a presentation, and submitting a final paper that details your analysis.

## Assignment 0

- ▶ We will be using R a *lot* in this course
- ▶ Download and install R (<http://www.r-project.org/>)
- ▶ You may want to use a nice IDE called RStudio  
(<http://www.rstudio.com/>)
- ▶ Install packages geoR, lattice, maps, maptools, spdep, spatstat, splancs

# Course Goals

- ▶ Through this course you will gain an understanding of what spatial data are, and how to work with these data.
- ▶ The key components of spatial analysis and statistics are:
  - ▶ **Visualization.** "A major pleasure in working with spatial data is their visualization. Maps are amongst the most compelling graphics, because the space they map is the space we think we live in, and maps can show things we cannot see otherwise".  
Bivand, Pebesma, Gomez-Rubio, Applied Spatial Data Analysis with R, Springer 2008.
  - ▶ **Exploration.** This involves looking for patterns in data such as clusters and the behavior of events that are close in space or very distant.
  - ▶ **Modeling.** We incorporate what we learn from data visualization and exploration into a formal statistical setting. This allows for estimation and inference.

# Historical examples of spatial analysis

## John Snow: Early spatial analysis

- ▶ In August 1854 there was a major Cholera outbreak in the Soho neighbourhood of London, UK. There were 127 cholera related deaths around the area.
- ▶ At the time, germ theory (microorganisms causing disease) was not really known.
- ▶ Dr. John Snow spoke to local residents and mapped where cholera cases occurred. As a result of his map, he was able to pinpoint the public water pump on Broad Street as the source of contaminated water causing the cholera outbreak.

# Historical examples of spatial analysis



# Historical examples of spatial analysis

## John Snow: Early spatial analysis

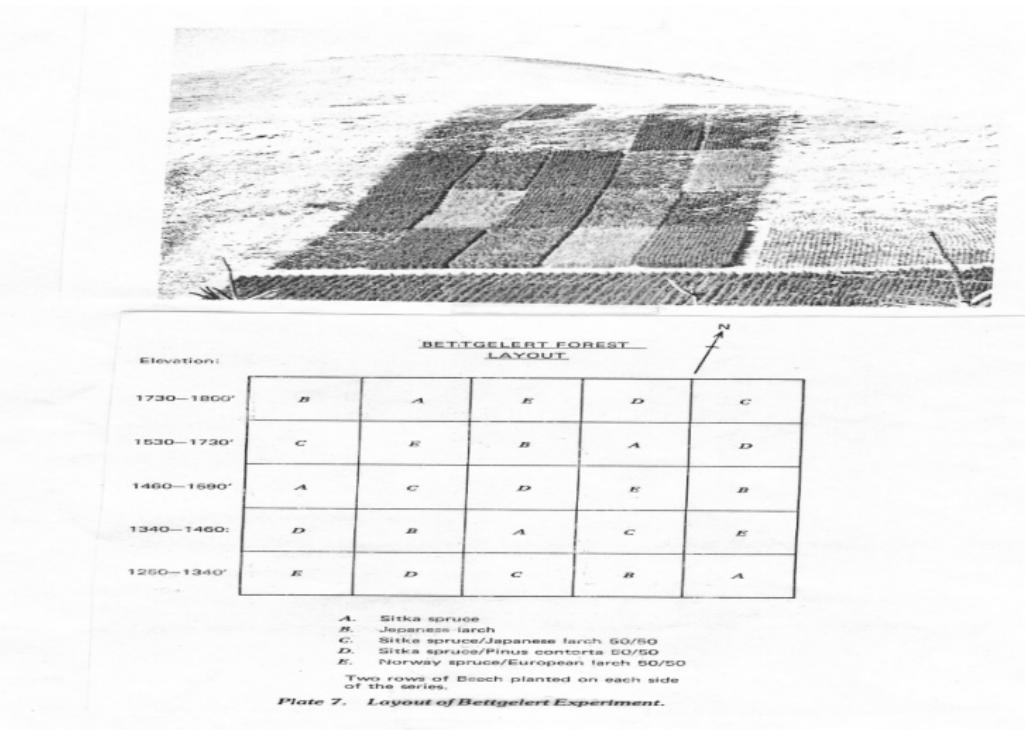
- ▶ Dr. Snow used statistics to find a relationship between water quality and cholera cases.
- ▶ He found that the waterworks company supplying water to Broad Street pump was taking water from the sewage polluted area of the Thames river.

# Historical examples of spatial analysis

## R.A. Fisher: Early spatial analysis

- ▶ R.A. Fisher was probably the first to recognize the implications of spatial dependence.
- ▶ In his work on design of experiments in agricultural science, he wrote (Fisher, 1935, p. 66):  
**After choosing the area we usually have no guidance beyond the widely verified fact that patches in close proximity are commonly more alike, as judged by the yield of crops, than those which are further apart.**
- ▶ Spatial variability, i.e. plot-to-plot variability, was largely due to physical properties of the soil and environmental properties of the field. He avoided the confounding of treatment effect with plot effect with the introduction of randomization
- ▶ Basically his solution was to eliminate spatial dependence by localizing into blocks.

# Historical examples of spatial analysis



## Spatial analyses: past and present

The historical examples were spatial analyses "by chance". Today, spatial data are seemingly everywhere. Not only do we see countless examples of spatial data in a variety of fields of research (public health, economics, sociology, earth sciences, etc.), we often find them used in news articles, apps and pop culture.

Let's first explain what is a spatial analysis and what are spatial data.

## What is spatial analysis?

- ▶ The quantification of phenomena referenced in space.
- ▶ The study of methods to describe and explain a process that operates in space based on a sample of observations taken at particular locations.
- ▶ Quantitative spatial analysis: Methods

## What is spatial analysis?

- ▶ The quantification of phenomena referenced in space.
- ▶ The study of methods to describe and explain a process that operates in space based on a sample of observations taken at particular locations.
- ▶ Quantitative spatial analysis: Methods
- ▶ **Visualization**  
Maps, graphical display

## What is spatial analysis?

- ▶ The quantification of phenomena referenced in space.
- ▶ The study of methods to describe and explain a process that operates in space based on a sample of observations taken at particular locations.
- ▶ Quantitative spatial analysis: Methods
- ▶ **Visualization**  
Maps, graphical display
- ▶ **Exploration**  
Tools to broadly look at spatial patterns

## What is spatial analysis?

- ▶ The quantification of phenomena referenced in space.
- ▶ The study of methods to describe and explain a process that operates in space based on a sample of observations taken at particular locations.
- ▶ Quantitative spatial analysis: Methods
- ▶ **Visualization**  
Maps, graphical display
- ▶ **Exploration**  
Tools to broadly look at spatial patterns
- ▶ **Modeling**  
Fitting models, testing hypothesis, formalizing spatial dependence

## What are spatial data?

- ▶ Data that are location specific and that vary in space
- ▶ Referenced by a spatial location,  $s$  where  $s = (x, y)$ ;  $x$  is longitude (easting) and  $y$  is latitude (northing). We typically have multiple spatial locations to examine, and they would be referenced as  $s_1, s_2, \dots, s_n = (x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ .
- ▶ Spatial locations may alternatively be referenced by an area such as a zip code, county, state. In such cases, we might use the centroid of the area as a reference point, and define the area and boundaries that encompass the area.
- ▶ In terms of thinking about how to approach the analysis of spatial data, the motto taken is:

## What are spatial data?

- ▶ Data that are location specific and that vary in space
- ▶ Referenced by a spatial location,  $s$  where  $s = (x, y)$ ;  $x$  is longitude (easting) and  $y$  is latitude (northing). We typically have multiple spatial locations to examine, and they would be referenced as  $s_1, s_2, \dots, s_n = (x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ .
- ▶ Spatial locations may alternatively be referenced by an area such as a zip code, county, state. In such cases, we might use the centroid of the area as a reference point, and define the area and boundaries that encompass the area.
- ▶ In terms of thinking about how to approach the analysis of spatial data, the motto taken is:
- ▶ **Data that are close together in space (time) are often more alike than those that are far apart.**

Often labeled as Tobler's first law of geography - "everything is related to everything else, but near things are more related than distant things".

# Types of Spatial Data

We can dissect the broad definition of spatial data into three sub-categories.

- ▶ **Point process data - Geostatistical data**
- ▶ **Areal data** (sometimes called aggregate data or lattice data)
- ▶ **Point pattern data**

There are specific statistical analyses for each of these types of spatial data. We will go into specifics of each.

# Types of Spatial Data

GEOSTATISTICAL DATA (POINT PROCESS)

# Geostatistical Data: An Everyday Example

Buy Rent Sell Mortgages Agent finder Advice Home design More

Address, Neighborhood, or ZIP  FOR SALE • \$500k+ • 1+ BEDS • HOME TYPE (1) • MORE •

SAVE SEARCH SAVED HOMES (0)

Real Estate 287 results.

Featured Price (High to Low) More

**605 N Linden Dr, Beverly Hills, CA**  
● HOUSE FOR SALE  
\$9,995,000  
5 bds • 5 ba • 6,246 sqft  
Open: Sat, 10am-12pm  
10 photos

**637 S Lucerne Blvd, Los Angeles...**  
● HOUSE FOR SALE  
\$5,388,000  
8 bds • 7 ba • 9,918 sqft  
Open: Sun, 2-5pm  
19 photos

**1525 Glen Oaks Blvd, Pasadena,...**  
● HOUSE FOR SALE  
\$5,200,000  
4 bds • 5 ba • 5,748 sqft  
Open: Sun, 1-4pm  
51 photos

**1623 Viewmont Dr, Los Angeles...**  
● HOUSE FOR SALE  
\$4,900,000  
4 bds • 5 ba • 3,723 sqft  
Open: Sun, 2-5pm  
50 photos

**1032 Hampton Rd, Arcadia, CA**  
● HOUSE FOR SALE  
\$4,588,000  
5 bds • 5.5 ba • 4,900 sqft  
Open: Sun, 2-4pm  
33 photos

# Geostatistical Data: Description

**Data that vary continuously over space, but is measured only at discrete locations**

Examples:

- ▶ housing prices in a metropolitan area
- ▶ field observations such as soil samples, air pollution measurements (environmental exposures)
- ▶ meteorological and climate data

**The common thread that links the data is a random process (also called stochastic process or random field)**

$$Z(\mathbf{s}) : \mathbf{s} \in D$$

where  $D$  is a domain in  $\mathbb{R}^d$  (d typically 2)

# Visualizing Geostatistical Data

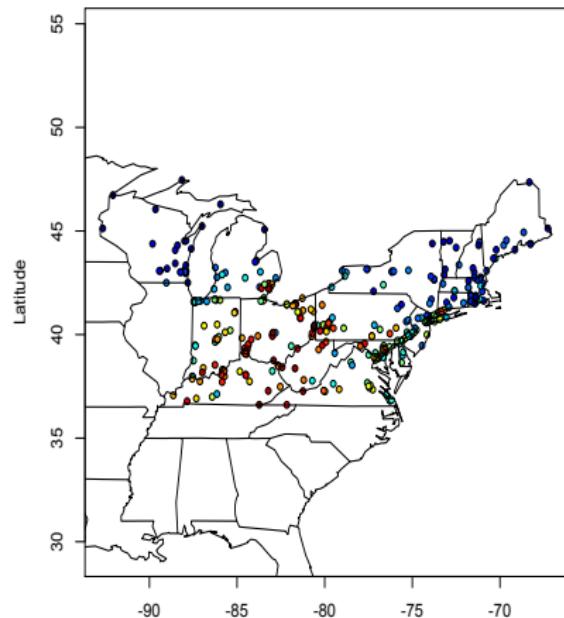
- ▶ The components of geostatistical data are the locations ( $s_i$ ), and the measurements at each location ( $z_s$ )
- ▶ The best way to visualize these data is to display on a map, and differentiate the values of the measurements of interest by colour or size.

Let's examine field observations of air pollution measurements in the northeast US.

# Visualizing Geostatistical Data

We display:

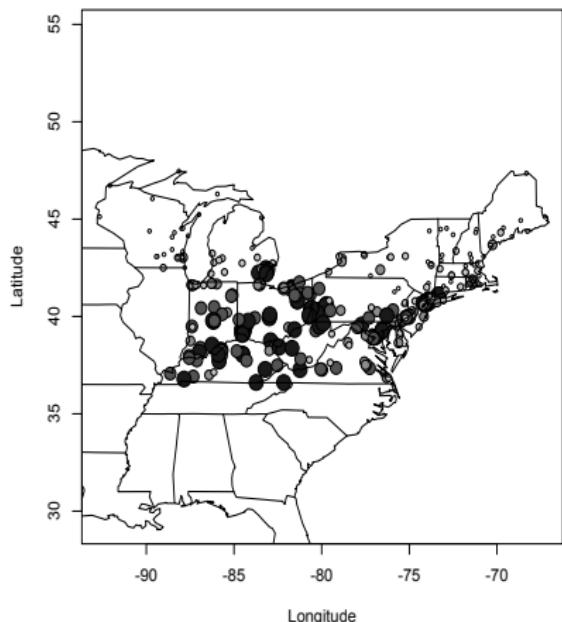
1. the points, which are air pollution monitors
2. the monthly average PM<sub>2.5</sub> concentration colour coded (a gradient from blue (low) to red (high))



# Visualizing Geostatistical Data

Alternatively we can display the points as gradients in size:

1. the points are air pollution monitors
  2. the monthly average PM<sub>2.5</sub> concentration where larger circles represent higher concentrations, smaller circles are lower concentrations
- \* Note the choice of colour and size gradient of the points can lead to different conclusions!



# Exploring and Modeling Geostatistical Data

## Goals of spatial statistics applied to geostatistical (point referenced) data

- ▶ Explore the spatial pattern in the observations. (Often called spatial "structure").
- ▶ Quantify the spatial pattern with a function.
- ▶ Model the spatial correlation/covariance in the observations.
- ▶ Make predictions at unobserved locations: interpolation, smoothing.

Additional considerations:

- ▶ Account for spatial structure in regression models.
- ▶ Test a null hypothesis of no spatial structure.

# Types of Spatial Data

AREAL DATA

# Areal Data: An Everyday Example

## Los Angeles Times | LOCAL

LOCAL U.S. WORLD BUSINESS SPORTS ENTERTAINMENT HEALTH STYLE TRAVEL OPINION SHOP

L.A. NOW POLITICS CRIME EDUCATION O.C. WESTSIDE NEIGHBORHOODS ENVIRONMENT DATA & MAPS LOCAL PLUS



MAPPING L.A.

### Neighborhoods

Search an L.A. County address

GO



#### Los Angeles County

Select one of 16 regions

- Angeles Forest
- Antelope Valley
- Central L.A.
- Eastside
- Harbor
- Northeast L.A.
- Northwest County
- Pomona Valley
- San Fernando Valley
- San Gabriel Valley
- Santa Monica Mountains
- South Bay
- South L.A.
- Southeast
- Verdugos
- Westside

## Areal Data: Description

- ▶ Data are associated with an area, so are aggregate in nature.
- ▶ Areal units tend to be irregular in shape (e.g. zip code, county) but can be regular grids (e.g. remote sensing data).
- ▶ Information collected in areal units may be census related, health related, environmental (satellite estimates of pollution, land cover).
- ▶ We want to determine spatial patterns of areal units within a region.
- ▶ Areal data (lattices) use neighbour relationships.
- ▶ Examples:
  - ▶ Median household income in Los Angeles neighborhoods
  - ▶ State-specific (or county-, census tract-, zip code-specific) election results
  - ▶ County hospital admission rates for influenza

# Visualizing Areal Data

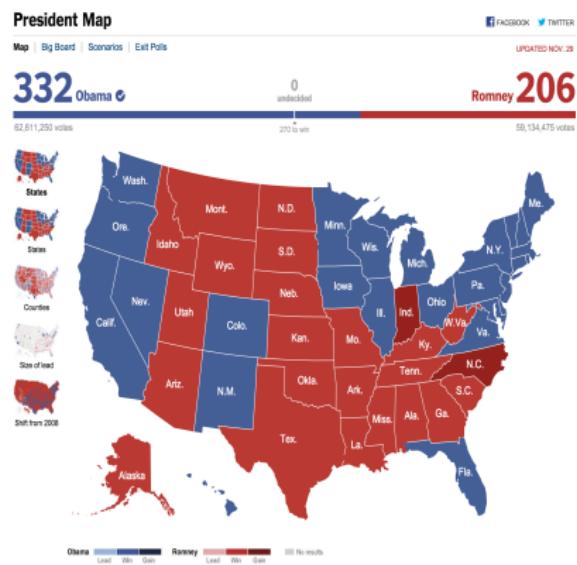
- ▶ Areal units are referenced as polygons.
- ▶ The centroids of the areal units may be useful for a spatial reference, in combination with the area of the polygon.
- ▶ The best way to visualize these data is to display as a map, differentiating the areal units by colour.

Let's examine the 2012 US presidential election results.

# Visualizing Areal Data: Example

We display:

1. the areal units (polygons), in this case, states.
2. the colour representing whether the state went to the Democrats (blue) or the Republicans (red).
3. a label of some sort, identifying the state.



# Visualizing Areal Data

What if we could break down the election results into counties?

- ▶ Smaller areas give us more spatial information.
- ▶ Visualizing the counties versus states should refine our view of the election results.
- ▶ Do we expect our map to show the same spatial pattern?

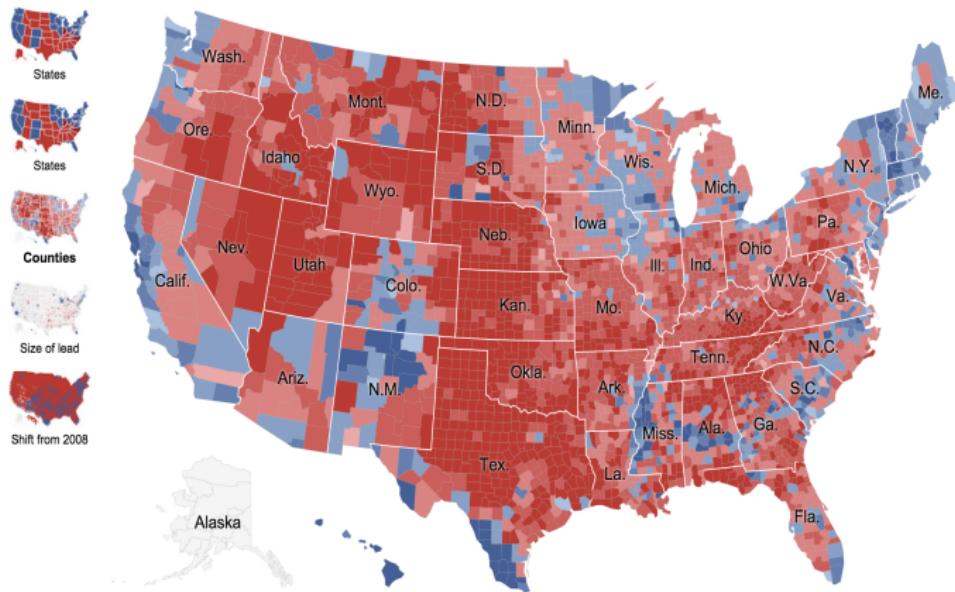
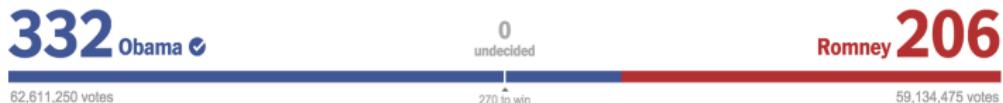
# Visualizing Areal Data: Example

## President Map

[FACEBOOK](#) [TWITTER](#)

[Map](#) | [Big Board](#) | [Scenarios](#) | [Exit Polls](#)

UPDATED NOV. 29



## Visualizing Areal Data

We notice that our visual conclusions can be different with different size areal units of the same thing. There often isn't anything we can do about it, other than to use statistical tools to analyze the data rather than just relying on visual conclusions. There are a couple of names of problems such as this that arise from areal data:

- ▶ The "modifiable areal unit problem" (MAUP)
- ▶ Ecological fallacy

Another problem we have is in the colour gradient of the areal units. Let's examine mortality rates in the US to illustrate this.

# Visualizing Areal Data

Crude birth rates by state based on equal-interval cut points

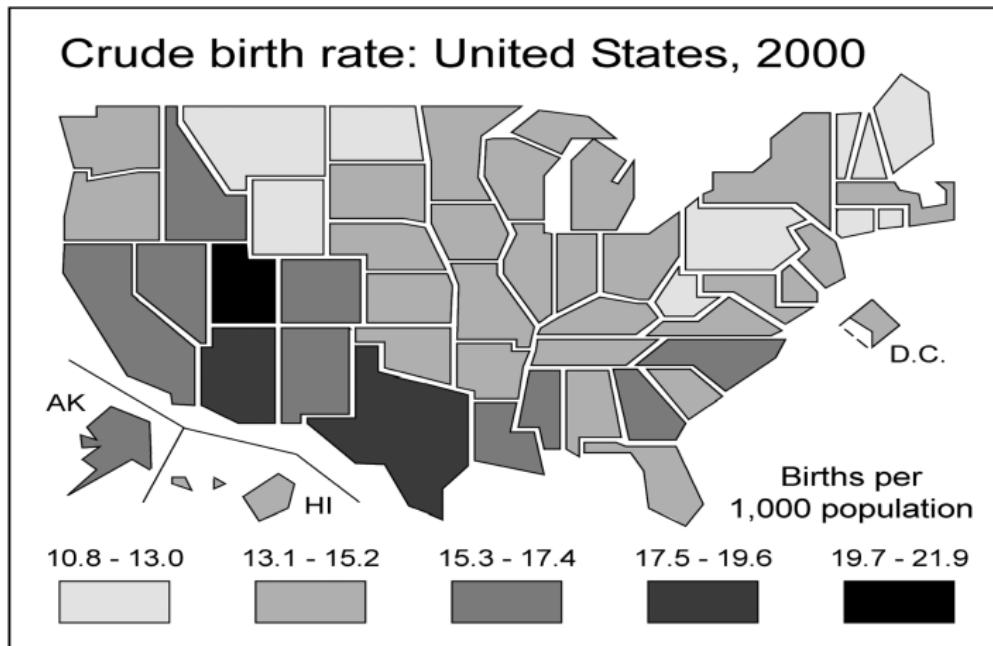


Figure: Monomier, N. Lying with Maps. Statistical Science 2005, 20(3) 215222.

# Visualizing Areal Data

Crude birth rates by state based on quantile cut points

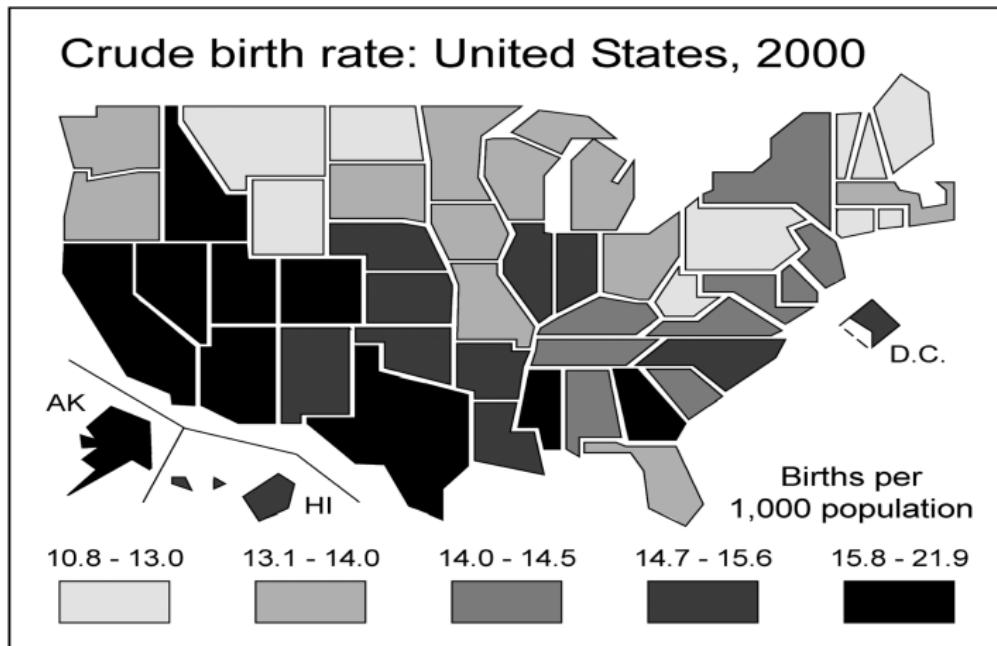


Figure: Monomier, N. Lying with Maps. Statistical Science 2005, 20(3) 215222.

# Exploring and Modeling Areal Data

## Is there a spatial pattern?

- ▶ Spatial pattern suggest that observations close to each other have more similar values than those far from each other.
- ▶ You might think that there is a pattern through visualization, but this is often subjective.
- ▶ Independent measurements will have no pattern, and would look completely random, but there may actually be an underlying pattern.

Goals of spatial statistics applied to areal data

- ▶ Understand the linkages between areal units and determine if there are areas that are closer to each other are more related than those that are farther apart.
- ▶ If there is a spatial pattern, *how strong is it?*

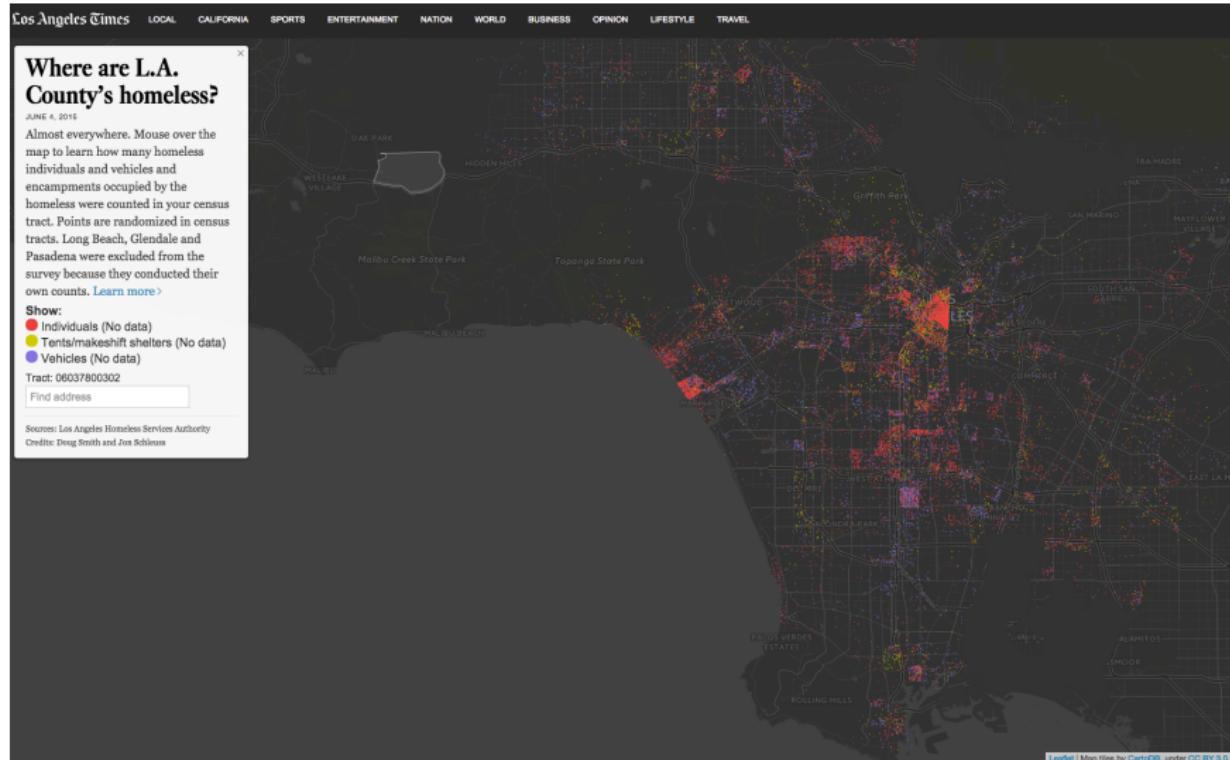
## Areal Data

The first reading assignment will be  
Mortimer, M (2005) Lying with maps. Statistical Science. 20(3)  
215-222.  
It is available on the course Blackboard. We'll discuss it next class.

# Types of Spatial Data

POINT PATTERN DATA

# Point Pattern Data: An Everyday Example



## Point Pattern Data: Description

- ▶ A spatial point process is a stochastic mechanism that generates events in 2D.
- ▶ Event is an observation (presence/absence), point is the location.
- ▶ Mapped point pattern: Events in a study area D have been recorded.
- ▶ Sampled point pattern: Events are recorded after taking samples in an area D.
- ▶ Examples
  - ▶ Locations of homeless in Los Angeles
  - ▶ Cases of malaria in Nairobi
  - ▶ Locations of a specific tree species in a forest

# Visualizing Point Pattern Data

- ▶ The components of point pattern data are the locations ( $s_i$ ).
- ▶ If there are different categories of a point pattern, such as with the homeless data, then these categories may be coloured separately.
- ▶ Often conclusions cannot be drawn from visual inspection alone

# Exploring and Modeling Point Pattern Data

Questions about point pattern data that we would like to answer are:

- ▶ Is there a regular pattern in the points?
- ▶ Is there clustering of the points?
- ▶ Can we define a point process that our events follow?

There may be additional features that we need to take into account:

- ▶ Is there an underlying population distribution from which events arise in a region?
- ▶ Are events clustering in areas of high population?
- ▶ If there are underlying features that could affect the presence/absence of a point, such as population density, we need to account for this.

# Exploring and Modeling Point Pattern Data

- ▶ Measure of intensity: mean number of events per unit area
- ▶ Are there differences between point process and a simple random process?
- ▶ Are points closer together than they would be by chance?
- ▶ Are the points more regularly spaced than they would be by chance?

# Exploring and Modeling Point Pattern Data

- ▶ Spatial location  $s$
- ▶ Presence/Absence modeled by  $Y$ ,  $Y(s) = 1$  if there is a case and  $Y(s) = 0$  otherwise
- ▶ Define a null hypothesis: no pattern (complete spatial randomness)
- ▶ Find a statistic to test whether the data is clustered, or regular.
- ▶ Model some spatial pattern and determine if our observed point pattern fits this model.

## Point Pattern Data

Let's look more closely at the interactive point pattern data of homelessness in Los Angeles.

<http://graphics.latimes.com/homeless-los-angeles-2015/>

## Types of spatial data: spatio-temporal

All three types of data we have described may be referenced in space and in time. That is, data that are location specific can have replicates in time:

- ▶ Each observation has a location, time and value
- ▶ Similar methods for analysis, with an added dimension
- ▶ Often encountered in environmental epidemiology:

**Geostatistical:** Relationship between daily air pollution measured at discrete locations in the US Northeast and hospital admissions

**Areal:** Examining birth rates from year to year in US states

**Point process:** Changes in spatial clustering of cholera in Kenya over time

## Important background concepts that play a role in this course

- ▶ Linear Regression:  $y_i = \beta_0 + \beta_1 x_i + \epsilon_i$
- ▶ The response variable  $y_i$  (predictand) is modeled as a linear function of one or more predictor variables  $x_i$  (covariates).
- ▶ The error,  $\epsilon_i$ , is associated with the response variable  $y_i$ ;
- ▶ The predictor variables are assumed to have no error (ignored, or  $x_i$ 's measured without error)

# Linear Regression: Matrix Notation

## Matrix Notation

- ▶ The general form of the linear model is

$$\mathbf{y} = \mathbf{X}\beta + \epsilon$$

- ▶ Where  $\mathbf{y}$  is the vector of responses (dependent variable) and  $\mathbf{X}$  is the "design matrix" of  $p$  explanatory variables (independent variables).

$$\mathbf{X} = \begin{bmatrix} 1 & x_{11} & \cdots & x_{p1} \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_{1n} & \cdots & x_{pn} \end{bmatrix}$$

$$\mathbf{y} = [y_1, \dots, y_n]$$

$$\epsilon = [\epsilon_1, \dots, \epsilon_n]$$

$$\beta = [b_0, b_1, \dots, b_n]$$

## Linear Regression: Least Squares

- ▶ To estimate  $\hat{\beta}$ , we minimize the squared error  $\sum_i \epsilon_i^2$
- ▶ Take the derivative, set to 0 to get the normal (or score) equations for  $\hat{\beta}$
- ▶ In matrix notation,  $\hat{\beta}_{OLS} = (X'X)^{-1}X'Y$

# Linear Regression

## Important Assumptions

- ▶ The residuals are identically and independently distributed (iid) from a normal distribution  $N \sim (0, \sigma^2)$
- ▶ In matrix notation, the residuals have mean zero and variance-covariance matrix  $\Sigma = \text{var}(\sigma) = \sigma^2 I$  where  $I$  is the identity matrix.
- ▶ The residuals have constant variance (homoscedastic)

## Linear Regression: Least Squares

- ▶ The OLS estimators of our regression parameters are unbiased (and the confidence intervals on the estimates are correct) when the model is correctly specified. Our covariates correctly specify the model and the previous assumptions are met.
- ▶ What if the assumptions fail?
- ▶ The variance-covariance matrix is not  $\sigma^2 I$
- ▶ There is covariance between errors, i.e.  $Cov(\sigma_i, \sigma_j)$
- ▶ Our variance-covariance matrix is  $\Sigma = var(\sigma) = \sigma^2 \mathbf{V}$

# Covariance

- ▶ Covariance measures linear associations.

$$Var(X) = Cov(X, X) = E[(X - E(X))^2] = E(X^2) - (E(X))^2$$

$$Cov(X, Y) = E[(X - E(X))(Y - E(Y))]$$

$$Corr(X, Y) = \frac{Cov(X, Y)}{\sqrt{Var(X)Var(Y)}}$$

## Linear Regression: Generalized Least Squares

- ▶ We need maximum likelihood methods to provide estimates for generating the score equations for  $\hat{\beta}$
- ▶ In matrix notation  $\hat{\beta}_{GLS} = (X'V^{-1}X)^{-1}X'V^{-1}Y$
- ▶ Familiar with MLE? We will go over in more detail when we talk about spatial regression.