

Spatial Data Analysis

Week 11: Point Pattern Data II

Meredith Franklin

Department of Statistical Sciences and School of the Environment

November 14, 2025

Outline

- ▶ Cluster point patterns
- ▶ Fitting point process models
- ▶ Machine learning for cluster detection: k-means, DBSCAN and HDBSCAN

Review: CSR, HPP, IPP

Definitions & properties

- ▶ CSR \equiv HPP on W : constant λ , independent counts.
- ▶ IPP: $\lambda(s)$ varies with s (covariates/trend); conditional independence.
- ▶ $N(A) \sim \text{Pois}(\mu(A))$; $\mu(A) = \lambda|A|$ (HPP) or $\int_A \lambda(s) ds$ (IPP).
- ▶ Likelihood $\propto \exp(-\int_W \lambda) \prod_i \lambda(x_i)$.
- ▶ CSR: $K(r) = \pi r^2$, $g(r) = 1$; $K > \pi r^2$ cluster, $K < \pi r^2$ inhibition.

Estimation & diagnostics

- ▶ HPP: $\hat{\lambda} = n/|W|$.
- ▶ IPP: $\log \lambda(s) = \beta_0 + \beta^\top z(s)$.
- ▶ Tests: quadrat; envelopes for K/L , G , H , etc.
- ▶ Inhomog.: K_{inhom} , G_{inhom} ; residual maps.

Cluster and Inhibition Processes

- ▶ Building block: Poisson process.
- ▶ Add-on: dependence between points.
- ▶ **Cox Process:** clustering arises because the intensity function is random; reflects a continuous latent field (doubly stochastic) and is more general.
- ▶ **Poisson Cluster Process (PCP):** parent–offspring mechanism.
- ▶ **Inhibition (SIP):** repulsion via interaction; regular spacing/competition.
- ▶ Example CSR baseline: $K(r) = \pi r^2$, $L(r) = r$
- ▶ Example Clustering: $K(r) > \pi r^2$ (or $L(r) > r$)
- ▶ Example Inhibition: $K(r) < \pi r^2$ ($L(r) < r$).

Cox Cluster Process

- ▶ A Cox process is a Poisson process driven by a random intensity surface $\Lambda(s)$.
- ▶ If two nearby locations s and $s+r$ tend to have simultaneously high intensity on that surface, then points are more likely to occur at both locations together—this is clustering.
- ▶ The “pair correlation” measures how often pairs at distance r occur relative to a Poisson baseline.
- ▶ **Cox process:** Poisson with *stochastic intensity* $\Lambda(s) \geq 0$. Conditional on Λ , points are IPP; clustering reflects

$$g(r) = 1 + \frac{\text{Cov}\{\Lambda(s), \Lambda(s+r)\}}{\lambda^2}, \quad \lambda = \mathbb{E}[\Lambda(s)].$$

- ▶ $\lambda = \mathbb{E}[\Lambda(s)]$ is the (mean) intensity.
- ▶ $g(r)$ is the pair correlation function (PCF).
- ▶ For stationary processes, $g(r) = \frac{\rho^{(2)}(r)}{\lambda^2}$, where $\rho^{(2)}(r)$ is the second-order product density at separation r . Note $\lambda(s)$ is first order intensity (could be written $\rho^{(1)}(s)$). Second order is joint intensity $\rho^{(2)}(s,t)$ expected number of ordered pairs with one point near s and one near t (for $s \neq t$).

Cox Cluster Process

- ▶ Why $g(r) = 1 + \text{Cov}\{\Lambda(s), \Lambda(s+r)\}/\lambda^2$ for Cox?
- ▶ Draw a random intensity field $\Lambda(s) \geq 0$; given Λ , points are IPP.
- ▶ Definitions: $\lambda = \mathbb{E}[\Lambda(s)]$; for stationary processes $g(r) = \rho^{(2)}(r)/\lambda^2$.
- ▶ Derivation:

$$\rho^{(2)}(r) = \mathbb{E}[\rho^{(2)}(s, s+r | \Lambda)] = \mathbb{E}[\Lambda(s)\Lambda(s+r)] = \lambda^2 + \text{Cov}\{\Lambda(s), \Lambda(s+r)\}.$$

$$\Rightarrow g(r) = \frac{\rho^{(2)}(r)}{\lambda^2} = 1 + \frac{\text{Cov}\{\Lambda(s), \Lambda(s+r)\}}{\lambda^2}.$$

- ▶ Interpretation: positive covariance at small $r \Rightarrow g(r) > 1$ (clustering); as covariance $\rightarrow 0$, $g(r) \rightarrow 1$ (Poisson-like).

Clustered Point Processes: Two Routes

- ▶ **Cox process:** Poisson with *stochastic intensity* $\Lambda(s) \geq 0$. Conditional on Λ , points are IPP; clustering reflects

$$g(r) = 1 + \frac{\text{Cov}\{\Lambda(s), \Lambda(s+r)\}}{\lambda^2}, \quad \lambda = \mathbb{E}[\Lambda(s)].$$

- ▶ **PCP (a special case of Cox):**

$$\Lambda(s) = \mu \sum_i h(s - s_i^{(P)}),$$

where parents P and their locations $\{s_i^{(P)}\}$ are Poisson with rate $\rho(s)$, N_i offspring/parent with $\mathbb{E}[N_i] = \mu$, and h is a kernel.

- ▶ **LGCP (another Cox subtype):** $\Lambda(s) = \exp(Z(s))$, Z Gaussian.

PCP: Definition

- ▶ A Poisson Cluster Process (PCP) generates events in $D \subset \mathbb{R}^2$ as **clusters** around **parent** points.
- ▶ Parent points form a Poisson process with intensity $\rho(s)$ (homogeneous: ρ constant).
- ▶ Each parent at location $s_i^{(P)}$ produces a random number of offspring N_i (iid; mean $\mathbb{E}[N_i] = \mu$).
- ▶ Offspring locations are $s_{ij} = s_i^{(P)} + u_{ij}$, where displacements u_{ij} are iid with pdf $h(u)$.
- ▶ Typical applications: seed/seedling patterns, insect larvae, galaxy clustering.

PCP: Mechanism

- ▶ **Parents:** HPP/IPP with rate ρ (or $\rho(s)$). Each parent i has iid offspring count N_i with $\mathbb{E}[N_i] = \mu$.
- ▶ **Displacements:** offspring offsets u_{ij} iid with pdf $h(\cdot)$; locations $s_{ij} = s_i^{(P)} + u_{ij}$.
- ▶ **Intensity field:**
$$\Lambda(s) = \mu \sum_i h(s - s_i^{(P)}), \quad \text{so } \lambda(s) = \mathbb{E}[\Lambda(s)] = \rho(s) \mu \text{ (homog.: } \lambda = \rho\mu).$$
- ▶ **PCP examples:** **Thomas** (Gaussian h), **Matérn cluster** (uniform-in-disk h).
Typical for parent–offspring mechanisms.
- ▶ **spatstat:** simulate `rThomas`, `rMatClust`.

PCP: Notation and First-Order Intensity

- ▶ Parent intensity: $\rho(s)$ on D .
- ▶ Cluster size: N_i with mean μ (e.g., Poisson, geometric, fixed).
- ▶ Displacement kernel: $h(u)$, a bivariate pdf ($\int_{\mathbb{R}^2} h(u) du = 1$).
- ▶ **Offspring intensity:** $\lambda(s) = \mu \rho(s)$ (homogeneous case: $\lambda = \mu \rho$).
- ▶ Isotropic kernels: $h(u) = \tilde{h}(\|u\|)$; e.g., Gaussian with variance σ^2 .

PCP: Displacement Kernels $h(u)$

- ▶ **Gaussian (Thomas process):**

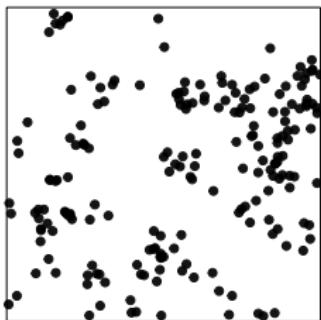
$$h(u) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{\|u\|^2}{2\sigma^2}\right).$$

- ▶ **Uniform-in-disk (Matérn cluster):** $h(u) = \frac{1}{\pi R^2}$ for $\|u\| \leq R$, else 0.
- ▶ Interpretation: σ^2 or R sets cluster spread; ρ controls number of clusters; μ controls cluster size.

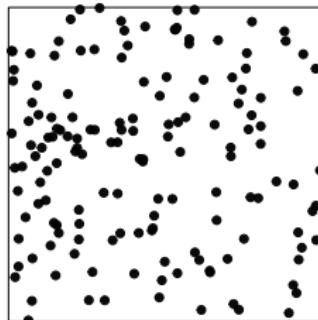
PCP Example: Thomas Process

Poisson Clustered Process

$$\rho = 50, \mu = 4, \sigma^2 = 0.001$$



$$\rho = 50, \mu = 4, \sigma^2 = 0.05$$



- D : unit square. Parents: $\rho = 50$. Cluster size: $E[S] = 4$.
- Left: smaller spread (e.g., Gaussian h with $\sigma^2 = 0.001$) \Rightarrow tight clusters.
- Right: larger spread (e.g., $\sigma^2 = 0.05$) \Rightarrow diffuse clusters.

PCP: Second-Order (Stationary/Isotropic)

Assume parent HPP with constant ρ ; isotropic kernel h ; cluster size N with $\mathbb{E}[N] = \mu$. Let $\lambda = \mu\rho$. For separation vector v with $r = \|v\|$:

$$\lambda_2(s, s+v) = \lambda^2 + \rho \mathbb{E}[N(N-1)] (h * h)(v), \quad (h * h)(v) = \int h(u) h(v-u) du.$$

Pair correlation: $g(r) = 1 + A (h * h)(r)$, $A = \frac{\mathbb{E}[N(N-1)]}{\mu^2 \rho}$.

$$K(r) = \pi r^2 + A \int_0^r 2\pi t (h * h)(t) dt.$$

Recall Poisson (CSR) baseline: $g(r) = 1$, $K(r) = \pi r^2$.

Nonstationary note: if parents have intensity $\rho(s)$, then $\lambda(s) = \mu \int h(s-u) \rho(u) du$ and use K_{inhom} , g_{inhom} .

Neyman–Scott: PCP Special Case

- ▶ Neyman Scott assumptions of **homogeneous** Poisson cluster process:
 - Parent events are realizations of a Poisson process with intensity ρ
 - Each parent i produces a random number of offspring N_i and the N_i are i.i.d.
 - The positions of offspring wrt the parent are i.i.d with bivariate pdf.
- ▶ **Neyman Scott (with Thomas Gaussian kernel, Poisson N):** $A = 1/\rho$, and

$$(h * h)(r) = \frac{1}{4\pi\sigma^2} \exp\left(-\frac{r^2}{4\sigma^2}\right), \quad g(r) = 1 + \frac{1}{\rho} \cdot \frac{1}{4\pi\sigma^2} \exp\left(-\frac{r^2}{4\sigma^2}\right).$$

- ▶ **Matérn (uniform disk radius R , Poisson N):** $A = 1/\rho$; $(h * h)(r)$ is the area-overlap kernel of two radius- R disks.
- ▶ Larger σ or $R \Rightarrow$ broader clusters; larger $\rho \Rightarrow$ more clusters but weaker $g(r) - 1$.

spatstat Pointers (PCP)

- ▶ **Simulate (PCP):** `rThomas(kappa=ρ, mu=μ, sigma=σ), rMatClust(kappa, mu, r).`
- ▶ **Fit (PCP):** `kppm(X ~ 1, clusters="Thomas") or "MatClust".`
- ▶ **Predict & simulate from fit:** `lam ← predict(fit, type="intensity"), simulate(fit, nsim=199).`
- ▶ **Diagnostics:** global envelopes for K, L, g ; intensity-weighted versions if trend suspected:
 - CSR envelopes: `envelope(X, Kest, nsim=199, global=TRUE).`
 - Inhomogeneous: `Kinhom(X, lambda=lam), pcfinhom(X, lambda=lam).`
 - Residual-style checks (from fitted Cox/cluster model): simulate from fit and compare summaries.

Cox Processes & Log-Gaussian Cox Process

- ▶ **Cox process:** Poisson with *stochastic intensity* $\Lambda(s) \geq 0$. Conditional on Λ , points are IPP; clustering reflects $\text{Cov}\{\Lambda(s), \Lambda(s+r)\}$ with

$$g(r) = 1 + \frac{\text{Cov}\{\Lambda(s), \Lambda(s+r)\}}{\lambda^2}, \quad \lambda = \mathbb{E}[\Lambda(s)].$$

- ▶ **LGCP (a Cox subtype):** $\Lambda(s) = \exp(Z(s))$ with Gaussian field Z . If Z is stationary with mean μ_Z , variance σ_Z^2 , correlation $\rho_Z(r)$:

$$\mathbb{E}[\Lambda] = \exp(\mu_Z + \frac{1}{2}\sigma_Z^2), \quad g(r) = \exp(\sigma_Z^2 \rho_Z(r)).$$

- ▶ **Examples (LGCP):** species occurrences (latent habitat), disease cases (latent risk), crime hotspots, rainfall/lightning fields.

spatstat Pointers (LGCP)

- ▶ **Simulate (LGCP):** `rLGCP()` with a covariance model (e.g., Matérn/exponential):
 - $X \leftarrow rLGCP(\text{model}=\text{"matern"}, \text{mu}=\eta, \text{var}=\sigma^2, \text{scale}=\alpha, \text{nu}=v, \text{win}=W)$.
 - η is the log-mean intensity (so $E[\Lambda] = \exp(\eta + \sigma^2/2)$).
- ▶ **Fit (LGCP):** either
 - `lgcp(X ~ trend, covmodel="matern", ...)` (direct LGCP fit), or
 - `kppm(X ~ trend, clusters="LGCP")` (version-dependent shortcut).
- ▶ **Predict intensity:** `lam ← predict(fit, type="intensity")` ⇒ `image(im)` for mapping and diagnostics.
- ▶ **Diagnostics (inhomogeneous):**
 - Pair correlation: `pcfinhom(X, lambda=lam)`; Poisson baseline is $g(r) = 1$.
 - K -function: `Kinhom(X, lambda=lam)`; compare against simulations from the fitted LGCP:
 - `envelope(X, Kinhom, nsim=199, simulate=expression(simulate(fit)), global=TRUE)`.
- ▶ **Trend + covariates:** include covariates in the trend (e.g., $X \sim z_1 + z_2$); supply rasters via `as.im()`.
- ▶ **Practical tips:** start with a simple trend, choose Matérn covariance (flexible v); check that estimated $g(r)$ shape matches model-implied $g(r) = \exp(\sigma_Z^2 \rho_Z(r))$.

LGCP vs PCP

- ▶ LGCP: $\Lambda(s) = \exp(Z(s))$, with Gaussian field Z .
- ▶ If Z stationary: mean μ_Z , var σ_Z^2 , corr $\rho_Z(r)$:

$$\mathbb{E}[\Lambda(s)] = \exp\left(\mu_Z + \frac{1}{2}\sigma_Z^2\right), \quad g(r) = \exp\left(\sigma_Z^2 \rho_Z(r)\right).$$

- ▶ PCP induces $g(r) = 1 + A(h * h)(r)$ (additive bump); LGCP induces multiplicative $g(r) = \exp(\cdot)$.

Fitting Point Process Models: Overview

- ▶ Data: observed points $\{x_1, \dots, x_n\}$ in region D .
- ▶ Choose a model class:
 - **IPP (trend only)**: $\lambda(s)$ varies with covariates/trend.
 - **Clustered Cox**:
 - **PCP / Neyman–Scott (Thomas, Matérn)**: parent → offspring mechanism.
 - **LGCP**: smooth latent field $\Lambda(s) = \exp(Z(s))$.
- ▶ Estimation routes:
 - **IPP**: log-linear likelihood via ppm.
 - **PCP/LGCP**: minimum contrast (least-squares in K or g) via kppm.
 - **Gibbs/interaction (e.g., Strauss)**: pseudolikelihood via ppm.
- ▶ Checking: $K/L/g$ envelopes (CSR or inhomogeneous), intensity-reweighted versions if trend present; simulate from fitted model.

Fit Inhomogeneous Poisson Process (IPP) via ppm

- ▶ Log-linear trend: $\log \lambda(s) = \beta_0 + \beta^\top z(s)$ (coordinates and/or covariates).
- ▶ spatstat: `ppm(X ~ z1 + z2 + x + y)`.
- ▶ Computation: maximizes the **log likelihood**

$$L(\theta | x) = \exp\left(-\int_D \lambda(u; \theta) du\right) \prod_{i=1}^n \lambda(x_i; \theta).$$

$$\ell(\theta | x) = \sum_{i=1}^n \log \lambda(x_i; \theta) - \int_D \lambda(u; \theta) du,$$

- ▶ Diagnostics: `Kinhom`, `pcfinhom`, residual maps; envelopes with `simulate=expression(rpoispp(`^lambda`))`.

PCP (Thomas, Matérn) via kppm (Minimum Contrast)

- ▶ Model: parents Poisson with rate ρ (or $\rho(s)$); each parent i has iid N_i with $\mathbb{E}[N_i] = \mu$; offsets $u_{ij} \sim h(\cdot)$.
- ▶ Homogeneous case: $\lambda = \mu\rho$; Thomas: Gaussian h (spread σ); Matérn: uniform-in-disk (R).
- ▶ **Fit in spatstat:**
 - `fitT <- kppm(X ~ 1, clusters="Thomas", method="mincontrast")`
 - `fitM <- kppm(X ~ 1, clusters="MatClust", method="mincontrast")`
- ▶ Statistic matched: by default $K(r)$ (or choose `statistic="pcf"`).
- ▶ Typical parameters returned: $\hat{\rho}$ (parent rate), $\hat{\mu}$ (mean cluster size), $\hat{\sigma}$ or \hat{R} (spread).
- ▶ Check: `simulate(fitT)`, `envelope(X, Kest, simulate=expression(simulate(fitT)), global=TRUE)`; also compare `pcf`.

LGCP via kppm (Minimum Contrast)

- ▶ LGCP: $\Lambda(s) = \exp(Z(s))$; Z Gaussian with mean μ_Z , variance σ_Z^2 , correlation (often Matérn) with range/scale α and smoothness v .
- ▶ **Fit in spatstat** (minimum contrast to K or g):
 - `fitL <- kppm(X ~ trend, clusters="LGCP", method="mincontrast")`
- ▶ Interpreting parameters: $\mathbb{E}[\Lambda] = \exp(\mu_Z + \frac{1}{2}\sigma_Z^2)$, $g(r) = \exp(\sigma_Z^2\rho_Z(r))$ (scale/smoothness shape g).
- ▶ Diagnostics: `pcf` or `pcfinhom` (if trend), envelopes via `simulate(fitL)`; intensity map: `predict(fitL, type="intensity")`.

Minimum Contrast: What kppm Optimizes

- ▶ Match summary function to theory by weighted least squares:

$$D(\theta) = \int_0^{r_{\max}} w(r) \{ \hat{S}(r)^c - S(r; \theta)^c \}^2 dr,$$

where $S \in \{K, g\}$, c is a power transform (variance stabiliser), $w(r)$ is a weight.

- ▶ Choices:
 - $S = K$: cumulative departures (broad clustering).
 - $S = g$: scale-specific departures (local range).
 - $c \approx 0.5$ (near-CSR), $c \approx 0.25$ (strong clustering); try a few to assess sensitivity.
 - r_{\max} : typically $\sim 1/3$ – $1/2$ of the shorter window side.
- ▶ Example (Thomas): $K(r; \rho, \sigma) = \pi r^2 + \frac{1}{\rho} (1 - e^{-r^2/(4\sigma^2)})$.
- ▶ Reference: Guan & Sherman (2007, JRSS B) for asymptotics; Baddeley et al. (2015) for practice.

Example: Fit Thomas PCP by K -contrast

Model: Thomas PCP with parameters $\theta = (\rho, \mu, \sigma)$ (ρ =parent intensity; μ =mean offspring/parent; σ =spread).

Intensity: $\lambda = \mu \rho$ (homogeneous).

Theoretical K (Thomas):

$$K_{\text{Th}}(r; \rho, \sigma) = \pi r^2 + \frac{1}{\rho} \left(1 - e^{-r^2/(4\sigma^2)} \right) \quad (\text{independent of } \mu).$$

Criterion (minimum contrast):

$$D(\rho, \sigma) = \int_0^{r_{\max}} \{ \hat{K}(r)^c - K_{\text{Th}}(r; \rho, \sigma)^c \}^2 dr,$$

choose c (e.g., 0.5; use 0.25 for strong clustering) and $r_{\max} \approx 1/3-1/2$ of the shorter window side.

Estimation: minimise $D(\rho, \sigma)$; then $\hat{\mu} = \hat{\lambda}/\hat{\rho}$ with $\hat{\lambda} = n/|D|$.

Note, we use ρ in the slides but it is 'kappa' in spatstat.

Testing & Diagnostics: What to Plot

- ▶ **CSR null:** `envelope(X, Kest)` and `envelope(X, pcf)` (global envelopes recommended).
- ▶ **Beyond trend:** fit IPP $\hat{\lambda}$, then `Kinhom(X, lambda=^ \lambda)`, `pcfinhom(X, lambda=^ \lambda)`; envelopes with `simulate=expression(rpoispp(^ \lambda))`.
- ▶ **Model-based:** compare to `simulate(fit)` for PCP/LGCP; plot pcf vs theoretical from `pcfmodel(fit)`.
- ▶ **Bandwidth for g :** use `bw.pcf(X)` to avoid over/under-smoothing.
- ▶ **Edge correction:** use `correction="iso"` (or translation) and restrict r to safe range.

Code Cheatsheet (spatstat)

- ▶ **IPP (trend)**: `fitI <- ppm(X ~ z1 + z2); lam <- predict(fitI, type="trend")`.
- ▶ **PCP (Thomas/Matérn)**: `fitT <- kppm(X ~ 1, clusters="Thomas"); fitM <- kppm(X ~ 1, clusters="MatClust")`.
- ▶ **LGCP**: `fitL <- kppm(X ~ trend, clusters="LGCP")`.
- ▶ **Envelopes**: `envelope(X, Kest, nsim=199, global=TRUE); envelope(X, pcfinhom, simulate=expression(rpoispp(lam))); envelope(X, Kest, simulate=expression(simulate(fitT)))`.
- ▶ **Pair correlation**: `pcf(X) (homog), pcfinhom(X, lambda=lam) (inhom); theoretical: pcfmodel(fitT)`.
- ▶ **Residuals/diagnostics**: `diagnose(fitI) or diagnose(fitT)`.

From Point Processes to Machine Learning

Spatial Clustering

- ▶ **Point process view:** a generative model for events in space (e.g., PCP, LGCP); clustering is inferred from $K(r)$, $g(r)$, parameters (ρ, μ, σ) , and model comparison.
- ▶ **ML clustering view:** an algorithmic partition of observed points into groups + noise, without specifying a generative process.
- ▶ **Bridge:** where $g(r) > 1$ at short ranges (excess close pairs), ML methods search for *locally dense* groups consistent with that signal.
- ▶ **Goal shift:** PP \Rightarrow estimate mechanisms & parameters; ML \Rightarrow discover structure, summarize shapes, detect anomalies.

Core Ingredients in ML Spatial Clustering

- ▶ **Inputs:** coordinates $s = (x, y)$ and optionally features $z(s)$.
- ▶ **Similarity:** a distance/affinity on points (Euclidean on projected coords; or geodesic for lon/lat).
- ▶ **Scale:** a neighborhood bandwidth (explicit or implicit) controlling what “dense” means.
- ▶ **Outputs:** cluster labels and possibly noise/outliers; sometimes soft membership or stability.
- ▶ **Validation:** internal indices (silhouette), stability under resampling, and spatial diagnostics (e.g., Moran’s I within clusters).
- ▶ **Good practice:** project to a planar CRS for meters, standardize heterogeneous features, set sensible spatial extents to reduce edge effects.

Method Families for Spatial Clustering

- ▶ **Centroid-based:** k -means
- ▶ **Density-based:** *DBSCAN* (radius ε , MinPts) finds high-density regions and labels noise; *HDBSCAN* builds a density hierarchy (no global ε) and selects stable clusters.
- ▶ **Hierarchical (agglomerative):** linkage on distances; dendrogram cut chooses granularity.
- ▶ **Model-based:** Gaussian Mixture Models (elliptical clusters; EM; choose k by AIC/BIC).
- ▶ **Graph/spectral:** build a k NN or radius graph and partition via eigenvectors; handles nonconvex shapes (parameter: neighborhood size).
- ▶ **Spatially constrained:** enforce adjacency/contiguity (e.g., region merging, spatial Ward, SKATER) when clusters must be connected areas.

From PP Clues to ML Choices

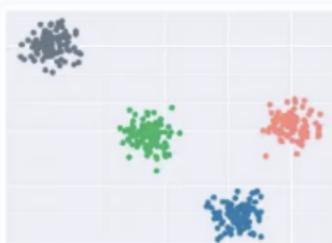
- ▶ **Use PP diagnostics to set scale:** if $g(r)$ peaks near r^* , start DBSCAN with $\varepsilon \approx r^*$; for varying density, prefer HDBSCAN.
- ▶ **Trend vs clustering:** if intensity varies strongly with s , de-trend or include features $z(s)$ before clustering; otherwise dense regions may reflect sampling effort, not structure.
- ▶ **Noise handling:** density-based methods provide an explicit noise label—useful when PP analysis suggested inhibition or sparsity between groups.
- ▶ **Report both:** ML clusters (maps, sizes, stability) *and* PP summaries ($K/L/g$) to justify the chosen scale and to compare alternative clusterings.

Density-Based Clustering (DBSCAN/HDBSCAN)

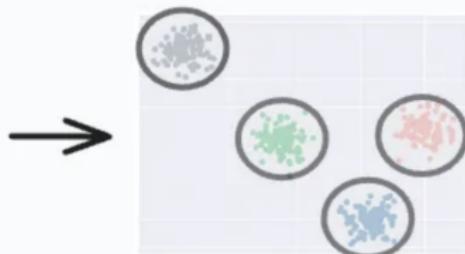
- ▶ **Idea:** clusters are regions of high point density separated by low-density gaps.
- ▶ **DBSCAN** (Ester et al., 1996): requires radius ε and MinPts; discovers arbitrary shapes and flags noise.
- ▶ **HDBSCAN**: hierarchical extension—*no global* ε ; builds a density tree and extracts stable clusters.
- ▶ Contrast to k -means: no centroids, no k chosen in advance, robust to outliers; shapes are non-convex.

k-means Spatial Clustering

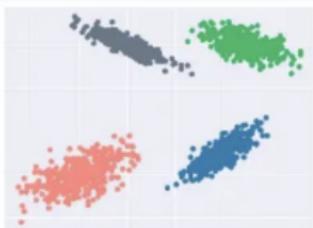
KMeans Clustering



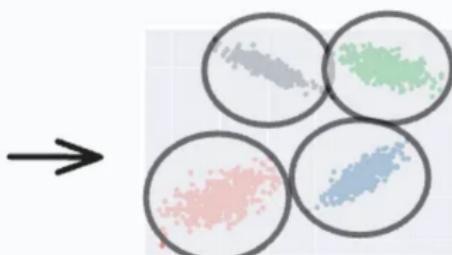
Data with
circular clusters



KMeans produces
round cluster

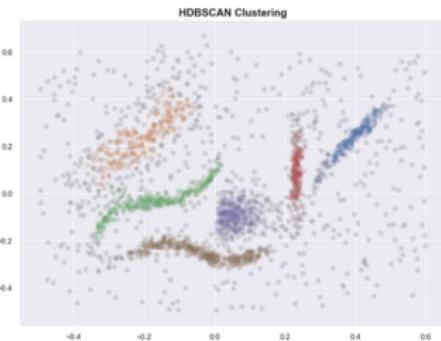
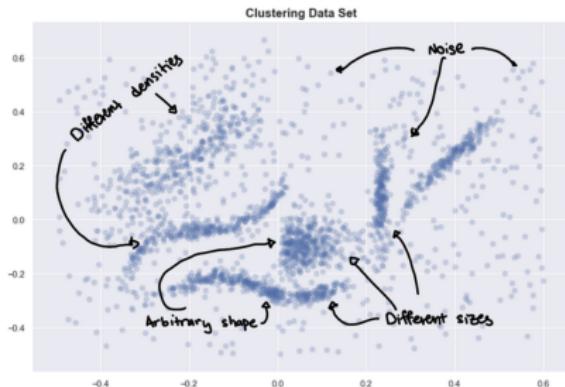


Data with
non-circular clusters



KMeans still
produces
round cluster

k-means vs DBSCAN



DBSCAN: Core Definitions

Given distance $d(\cdot, \cdot)$, radius $\varepsilon > 0$, and MinPts:

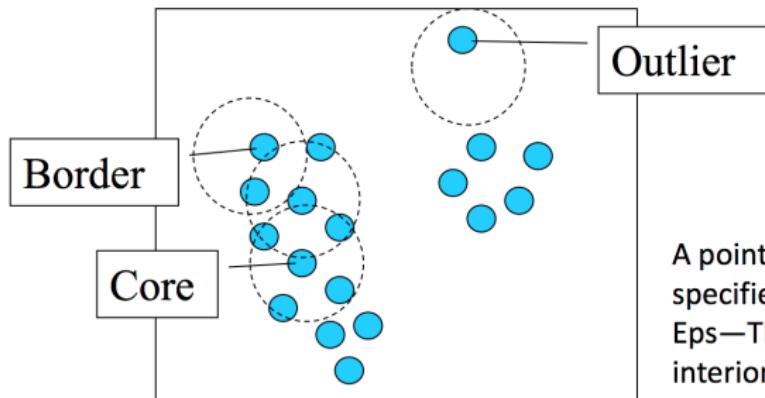
- ▶ **ε -neighborhood:** $N_\varepsilon(x) = \{y : d(x, y) \leq \varepsilon\}$. The count of neighbors within radius ε .
- ▶ **Core point:** $|N_\varepsilon(x)| \geq \text{MinPts}$. This is how a core point is determined.
 $|N_\varepsilon(x)| < \text{MinPts}$ says x does not have enough neighbors to be a core point (it will be noise or a border of a different core point).
- ▶ **Border point:** not core, but within ε of a core.
- ▶ **Noise:** neither core nor border.
- ▶ **Directly density-reachable:** $y \in N_\varepsilon(x)$ with x core.
- ▶ **Density-reachable:** connected by a chain of directly density-reachable steps.
- ▶ **Cluster:** maximal set of points density-connected to each other.

Density Based Clustering

The ε defines the neighborhood, where points within ε radius from another point are considered part of a cluster as long as it fulfills that there are a certain number of MinPts.



Density Based Clustering



$$\epsilon = 1 \text{ unit}, \text{MinPts} = 5$$

Given ϵ and MinPts , categorize the objects into three exclusive groups.

A point is a **core point** if it has more than a specified number of points (MinPts) within ϵ —These are points that are at the interior of a cluster.

A **border point** has fewer than MinPts within ϵ , but is in the neighborhood of a core point.

A **noise point** is any point that is not a core point nor a border point.

DBSCAN: Algorithm and Parameters

Algorithm (sketch):

1. Visit an unassigned point x ; if $|N_\varepsilon(x)| < \text{MinPts}$, mark as noise (may be upgraded to border).
2. If core, start a new cluster and iteratively absorb all points density-reachable from x .

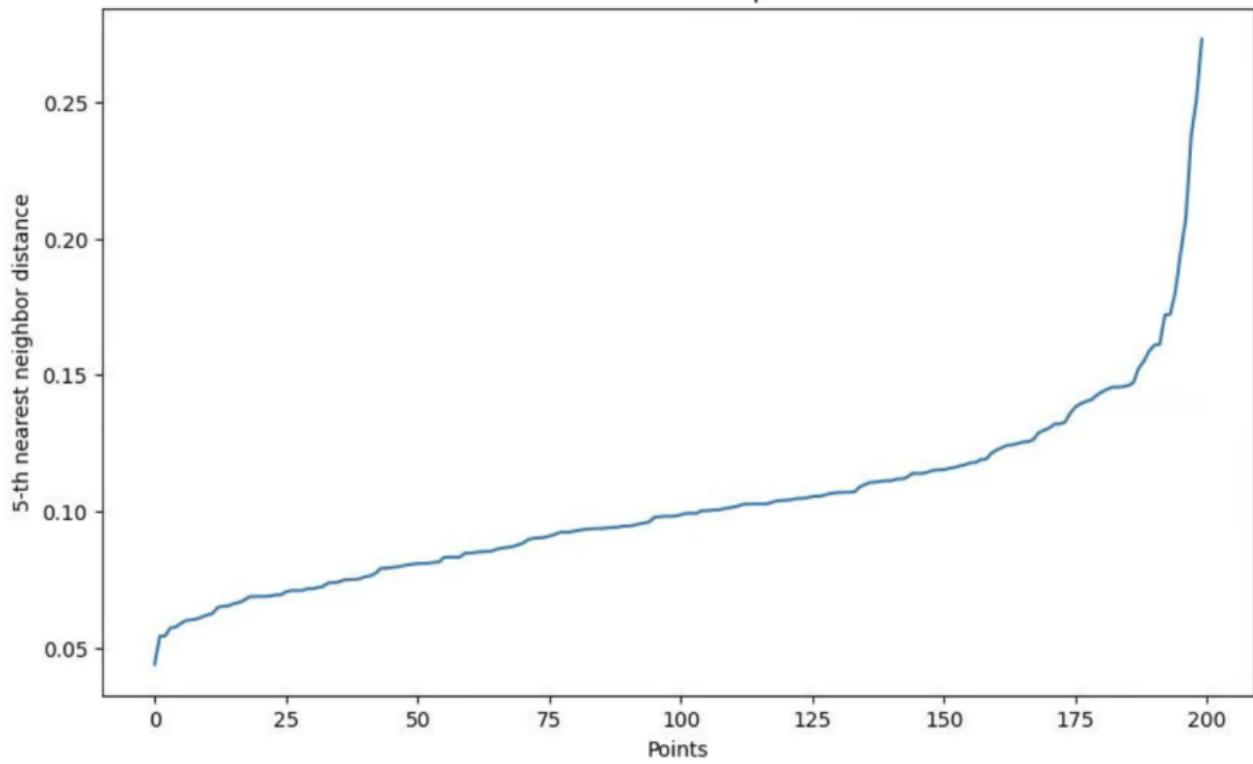
Parameters:

- ▶ MinPts : rule-of-thumb $\geq d+1$ in d dimensions (e.g., 4–5 in 2D), larger in noisy data.
- ▶ ε : pick via the **k -distance plot** (with $k = \text{MinPts}$); choose the “elbow.”

Complexity: $O(n \log n)$ with spatial index (kd/ball tree); $O(n^2)$ naive.

DBSCAN k-distance

K-distance Graph

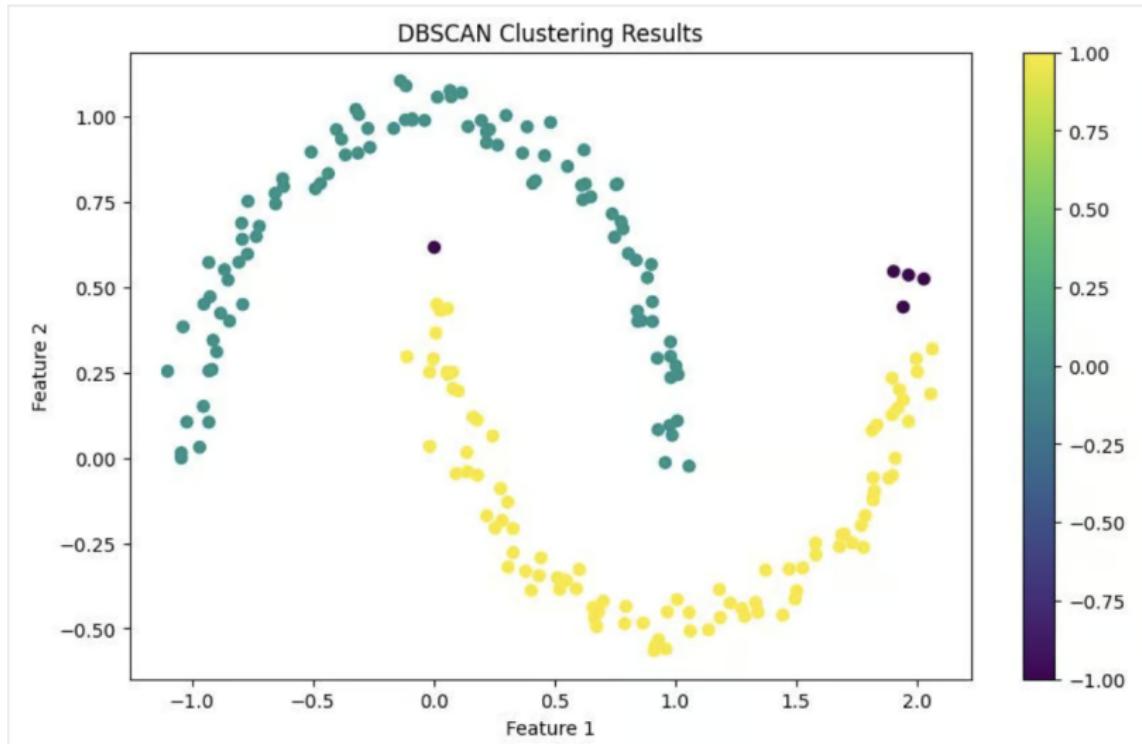


pick $\varepsilon = 0.15$ based on this plot

DBSCAN: Practical Considerations

- ▶ **Distance:** choose a metric appropriate to the data; for latitude/longitude either project (e.g., UTM) or use geodesic distances.
- ▶ **Scaling:** standardize heterogeneous features before Euclidean distance.
- ▶ **Strengths:** arbitrary shapes, automatic noise detection, no k .
- ▶ **Limitations:** single global $(\varepsilon, \text{MinPts})$ struggles with *varying density*; degraded in high dimension.

DBSCAN output



DBSCAN and HDBSCAN

- ▶ Here is a good visualization of DBSCAN results:
<https://www.naftaliharris.com/blog/visualizing-dbscan-clustering/>
- ▶ A better version in Hierarchical DBSCAN (HDBSCAN), where we do not need to set the ε neighborhood. With HDBSCAN we do set the minimum number of points that we wish to have in a cluster (2 to n). A larger MinPts means larger clusters.
 - HDBSCAN uses the concept of mutual reachability, where we look at distances that connect all of the points
 - A tree (hierarchy) is formed based on these distances

HDBSCAN: Why and How

Motivation: avoid a global ε ; handle clusters of *varying density*.

Key constructs:

- ▶ **Core distance** (order k): $c_k(x) = \text{distance to the } k\text{-th nearest neighbor}$ ($k = \text{MinPts}$).
- ▶ **Mutual reachability distance**:

$$m_k(x, y) = \max\{c_k(x), c_k(y), d(x, y)\}.$$

- ▶ Build an MST (minimum spanning tree) on the graph weighted by $m_k(\cdot, \cdot)$; varying the density level induces a **cluster hierarchy**.

Minimum Spanning Tree (MST)

Definition. For a connected, undirected, weighted graph $G = (V, E, w)$, a minimum spanning tree is a subset of edges $T \subseteq E$ that (i) connects all vertices (spans), (ii) has no cycles (is a tree), and (iii) has minimum total weight $\sum_{e \in T} w(e)$ among all spanning trees.

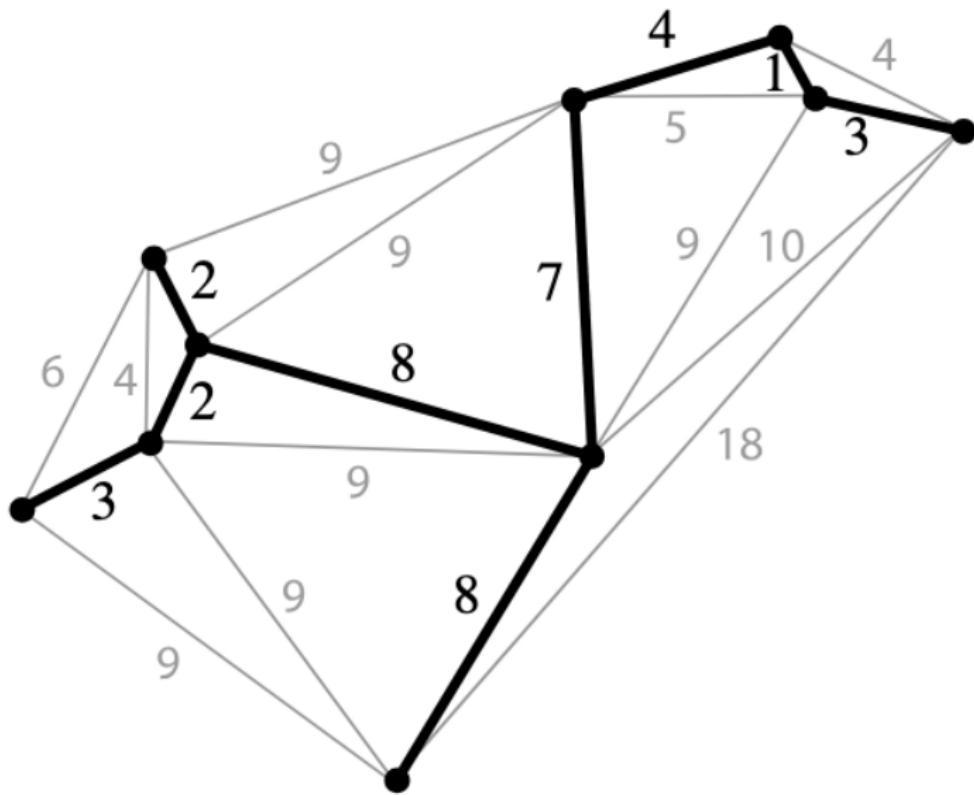
Key facts

- ▶ $|T| = |V| - 1$ edges; acyclic; connects all nodes.
- ▶ **Uniqueness:** if all edge weights are distinct, the MST is unique.
- ▶ **Cut property:** for any cut of V , the minimum-weight edge crossing the cut belongs to *some* MST.
- ▶ **Cycle property:** in any cycle, the maximum-weight edge cannot belong to *any* MST.
- ▶ **Algorithms:** Kruskal's ($O(E \log E)$), Prim's ($O(E \log V)$), Boruvka's (parallel-friendly).

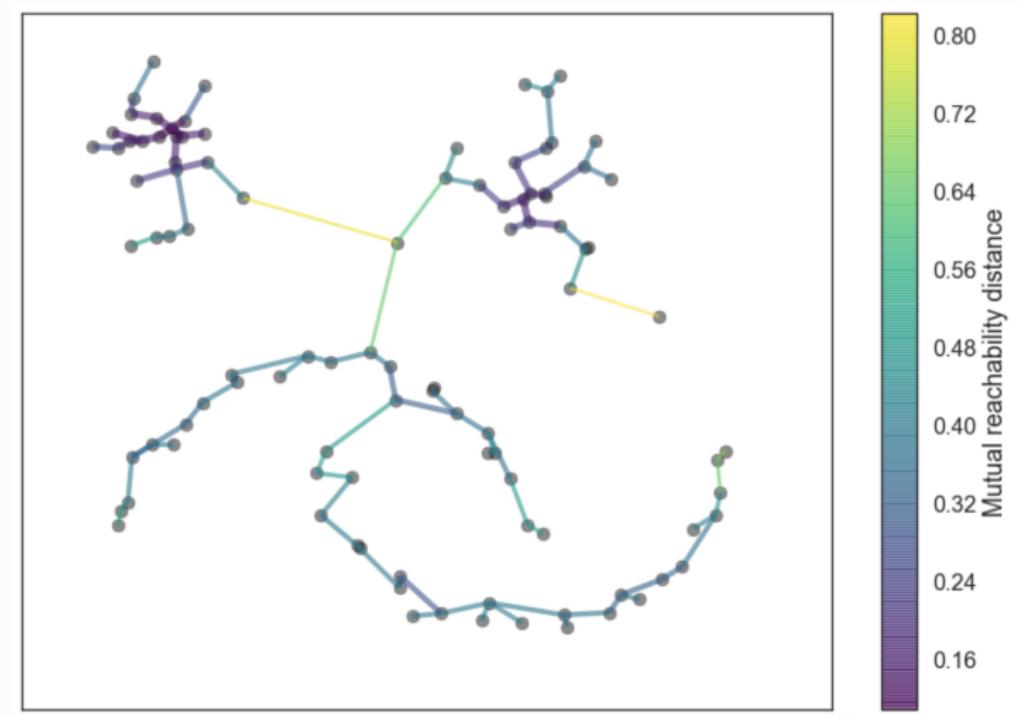
Why for clustering (DBSCAN/HDBSCAN)

- ▶ Single-linkage clustering = building an MST, then cutting long edges.
- ▶ HDBSCAN builds the MST of mutual reachability distances, then extracts stable clusters from the hierarchy.
- ▶ For large spatial datasets, sparsify with a k NN graph before MST.

Minimum Spanning Tree



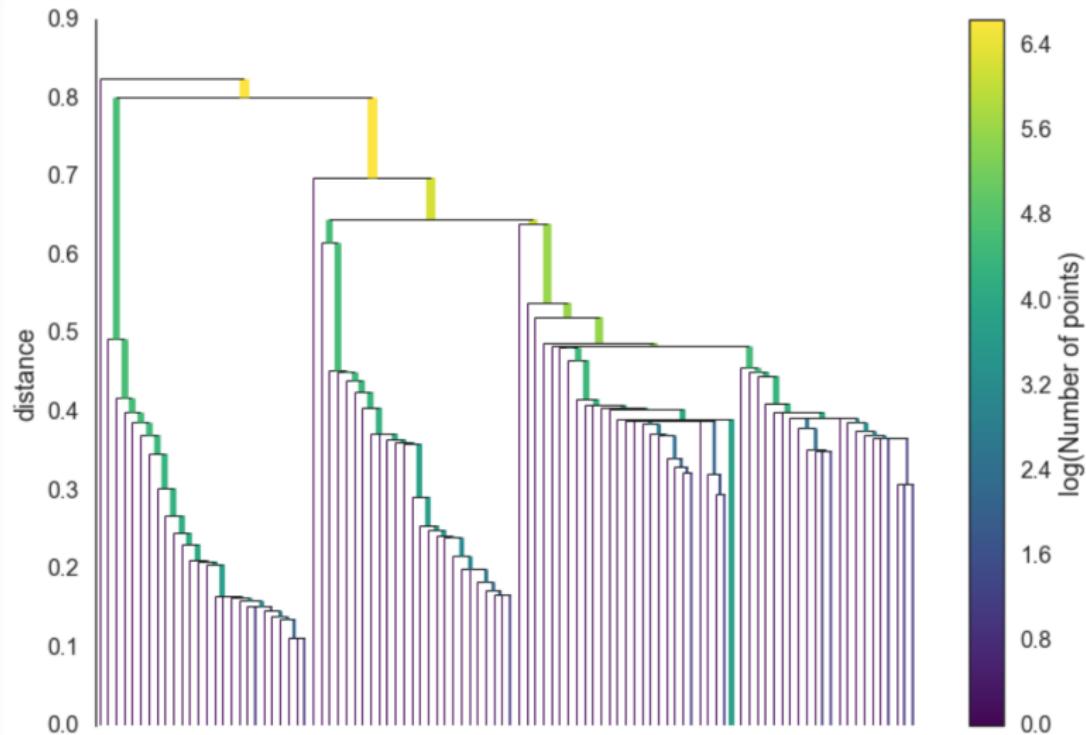
HDBSCAN Minimum Spanning Tree



HDBSCAN: From MST to Hierarchy

- ▶ **Input:** MST of *mutual reachability* distances $m_k(u, v)$ on the points.
- ▶ **Goal:** build a hierarchy (single-linkage dendrogram) of connected components.
- ▶ **Procedure (sweep on edge weights):**
 - Sort MST edges by distance (ascending).
 - Initialize each point as its own cluster (use a union–find/DSU structure).
 - For each edge (u, v, w) in order:
 - Find current clusters $C(u)$ and $C(v)$.
 - If $C(u) \neq C(v)$, merge them; record a dendrogram merge at level w .
 - Continue until all points are connected (one component).
- ▶ **Output:** a dendrogram where branch heights correspond to mutual reachability distance.
- ▶ **Notes:**
 - Cutting the dendrogram at a fixed level (fixed ε) reproduces DBSCAN.
 - HDBSCAN then **condenses** the tree using `min_cluster_size` and selects **stable** clusters.

HDBSCAN tree

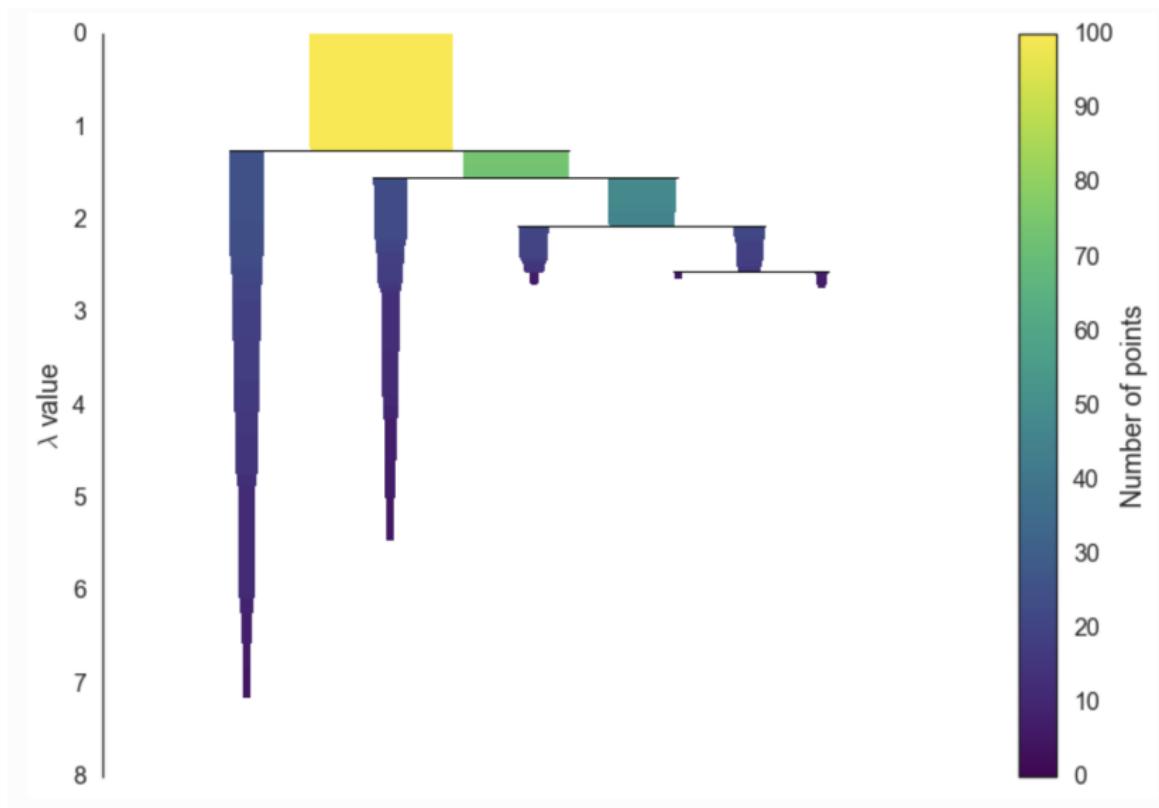


If we cut the tree at a given distance (ε) that is giving us the same estimate as DBSCAN

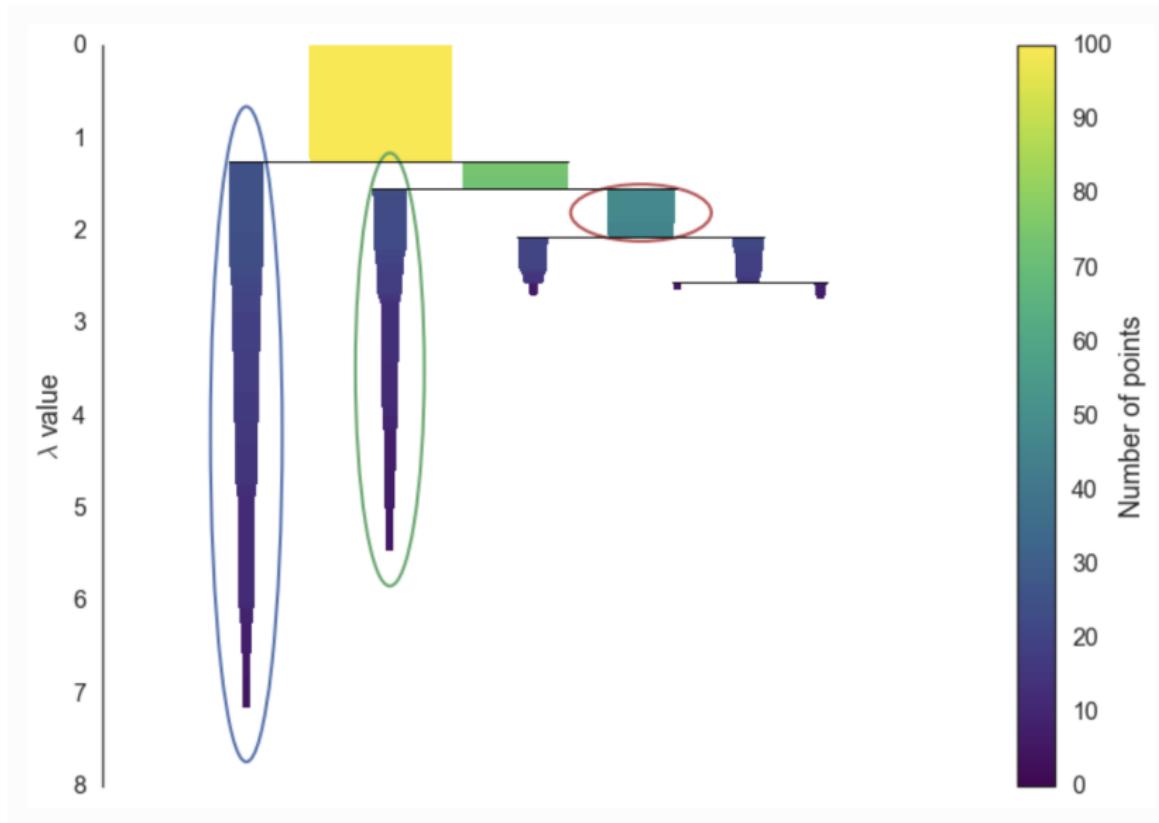
HDBSCAN: Condensed Tree and Cluster Selection

- ▶ Convert the hierarchy to a **condensed tree** using a `min_cluster_size` (often $\geq \text{MinPts}$).
- ▶ At each split, discard branches with fewer than `min_cluster_size` points.
- ▶ **Select clusters** by maximizing **stability** (excess-of-mass): favor branches that *persist* over wide density ranges.
- ▶ Output: hard labels (clusters + noise), optional soft memberships and cluster *persistence* scores.

HDBSCAN condensed tree



HDBSCAN condensed tree



HDBSCAN: Cluster Stability

- ▶ Use density level $\lambda \equiv 1/(\text{mutual reachability distance})$.
- ▶ For each cluster C in the *condensed tree*:
 - **Birth** $\lambda_{\text{birth}}(C)$: first λ where $|C| \geq \text{min_cluster_size}$.
 - **Death** $\lambda_{\text{death}}(C)$: λ where C splits or disappears.
 - For each point $p \in C$, let λ_p be the λ when p leaves C (to a child or to noise); if it never leaves before C dies, set $\lambda_p = \lambda_{\text{death}}(C)$.
- ▶ **Stability** (a.k.a. excess of mass):

$$\text{Stability}(C) = \sum_{p \in C} (\lambda_p - \lambda_{\text{birth}}(C)).$$

- ▶ **Selection rule** (dynamic programming on the tree): choose children if their total stability exceeds the parent's; otherwise keep the parent. This yields a non-overlapping set of *stable* clusters.

HDBSCAN: Membership Probability

- ▶ After selecting clusters, assign a **probability** to each point p for its cluster C based on how long p *persists* in C above birth:

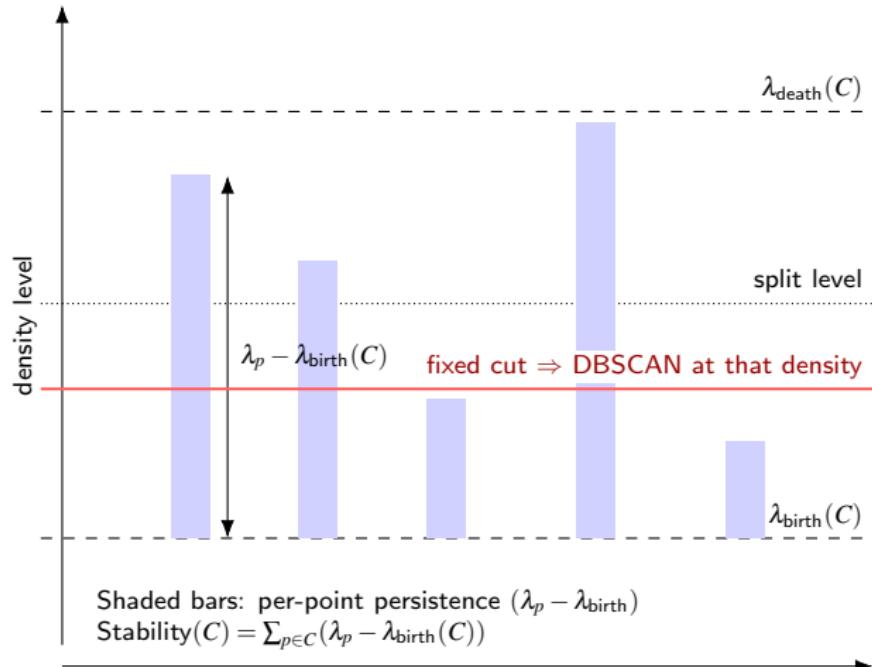
$$\Pr(p \in C) = \frac{\lambda_p - \lambda_{\text{birth}}(C)}{\lambda_{\text{death}}(C) - \lambda_{\text{birth}}(C)} \in [0, 1],$$

clipped to $[0, 1]$. (Intuition: 1 if p stays with C for its whole lifetime; smaller if p peels off early.)

- ▶ Implementation details:
 - The Python `hdbscan` package computes equivalent *persistence-based* scores; formulas are normalized so the most persistent points get probability near 1.
 - Points never in any selected cluster get probability 0 (noise).
- ▶ Practical readout: report **cluster stability** (one score per cluster) and **membership probabilities** (one per point) to summarize confidence.

HDBSCAN: Cluster Stability Schematic

$$\lambda = 1/(\text{mutual reachability})$$



DBSCAN vs HDBSCAN

	DBSCAN	HDBSCAN
Parameters	ε , MinPts	MinPts, <code>min_cluster_size</code> (no global ε)
Density	Single global threshold	Varying density via hierarchy
Noise	Explicit label	Explicit label; soft membership
Shapes	Arbitrary	Arbitrary
When	Uniform-density clusters	Mixed-density clusters; parameter-robust

R Workflow: DBSCAN and HDBSCAN

- ▶ **Data prep** (project & extract coords):

```
library(sf); library(dbSCAN)
pts <- st_transform(pts_sf, 32610) # project (ex: UTM zone 10N)
X   <- st_coordinates(pts)       # n x 2 matrix
```

- ▶ **Pick ε via k -NN distance plot** (use $k = \text{MinPts}$):

```
k <- 5
kNNdistplot(X, k = k); abline(h = 120, lty = 2) # knee near 120 m -> eps
```

- ▶ **Run DBSCAN:**

```
mod_db <- dbSCAN(X, eps = 120, minPts = k)
table(mod_db$cluster) # 0 = noise
```

- ▶ **Run HDBSCAN (no eps):**

```
mod_hd <- hdbscan(X, minPts = 5, minClusterSize = 12)
mod_hd$cluster_scores # stability/persistence
```

- ▶ **Plot:** color by cluster; mark noise separately.

Tip: scale non-spatial features before combining with coordinates.